

# Transformers applied to numerical simulator modelling

Max Cohen

March 09 2020

# Outline

- 1 Neural networks
- 2 Deep learning for Time Series
- 3 Transformer
- 4 Application on Oze data

# Definition

Consider a problem where we want to predict  $Y$  given an input  $X$ .  
 A neural network is a function  $F_\theta$  with parameters  $\theta$  called  
 "weights" that computes:

$$F_\theta(X) = \mathbb{P}_\theta(Y|X)$$

in a **classification setting** ( $Y$  lies in a discrete set) or


$$F_\theta(X) = \hat{Y}$$

in a **regression setting** ( $Y$  lies in a subset of  $\mathbb{R}^d$ ).

Weights are estimated using a dataset  $(X^{(i)}, Y^{(i)})_{1 \leq i \leq m}$ , by  
 minimizing a loss function  $L : (x, y) \mapsto \mathbb{R}^+$ :

$$\theta \mapsto \frac{1}{n} \sum_{i=1}^n L(F_\theta(X^{(i)}), Y^{(i)}).$$

# Fully connected neuron



A **linear transform** followed by a **nonlinear activation**:

$$F_{\theta}(x) = \sigma(\theta \cdot x + b), \quad \theta \in \mathbb{R}^n, b \in \mathbb{R},$$

where  $\sigma$  is a non-linearity, for instance a sigmoid function  $\sigma : x \mapsto \frac{1}{1+e^{-x}}$  or a ReLU  $\sigma : x \mapsto \max(0, x)$ .

# Bigger architectures





Figure: Fully connected network



Figure: Inception network

# Problem definition



We consider a **many to many regression problem**, for instance predicting the **indoor temperature** given weather data, or solving a **Natural Language Processing (NLP)** task.

$$X : (T, n)$$

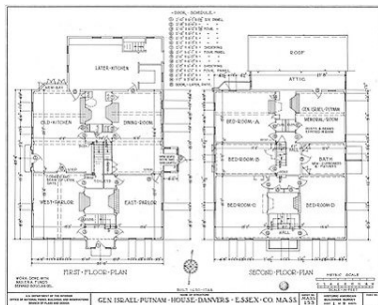
$$Y : (T, d)$$

## Example: consumptions data

→ **Predict consumptions and temperatures in a building**

- **Input**  $X = (X_1, \dots, X_T)$  :
  - **Time series** such as humidity, outside temperatures.
  - Received **hourly**.
- **Output**  $Y = (Y_1, \dots, Y_T)$ :
  - Inside temperatures **in several spots**.
  - Consumptions (**cooling, heating, ventilation**).
- **Controlled system** :
  - The time series also depend on a building management system (**not mentionned today**).

## Example: consumptions data



Model based on a **schematic view of the building** (geometry, etc.) and a **simulator for transient systems**.



## Example: consumptions data

Physics based approaches to propose simulators are widespread in the most complex situations.

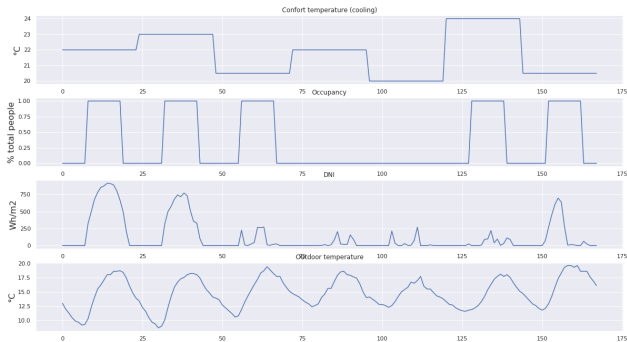
The main drawbacks of highly sophisticated softwares to simulate the behavior of transient systems: significant computational time required to train such models and integration of many noisy sources.

- 1 Millions of data obtained from a a physics based approach / software.
- 2 (STEP 1) Design a new flexible model (deep learning for time series) and train model parameters during an initial phase with these measurements.


## Example: consumptions data

- 1 Millions of data obtained from a **a physics based approach / software**.
- 2 **(STEP 1)** **Design a new flexible model** (deep learning for time series) and **train model parameters** during an initial phase with these measurements.
- 3 **(STEP 2a)** Analysis of **the statistical predictive efficiency** of the model. Uncertainty quantification in the prediction using sensor data.
- 4 **(STEP 2b)** Tuning of commands to ensure **energy savings** with **no renovation works**.

## Example: input data



## Example: output data



# Recurrent Neural Networks





Figure: RNN architecture

Until recently, **recurrent neural networks (RNN)** had been established as the go-to architecture for sequence modelling. They easily scale to long sequences, with a complexity in  $\mathcal{O}(T \cdot n^2)$ .

## Traditional architecture




At time  $t$ , a **hidden state of the network** is computed as follows:

$$h_t = \sigma_h(W_x x_t + W_h h_{t-1} + b_h),$$

where  $\sigma_h$  is a nonlinear activation function.

$W_x$ ,  $b_h$  and  $W_h$  are the **unknown parameters** of the state update.


## Traditional architecture



At time  $t$ , the **hidden state of the network** is used to predict  $Y_t$  as follows:

$$\hat{y}_t = F_{\theta}(x_t, h_{t-1}) = \sigma_y(W_y \cdot h_t + b_y).$$

# Long Short Term Memory (LSTM)



The LSTM cell is a more complex recurrent neural network. It contains three gates, **input**, **forget**, **output gates** and a **memory cell**. The **gates at time  $t$**  are:


$$\Gamma_i = \sigma(W_i[h_{t-1}, x_t] + b_i)$$

$$\Gamma_f = \sigma(W_f[h_{t-1}, x_t] + b_f)$$

$$\Gamma_o = \sigma(W_o[h_{t-1}, x_t] + b_o)$$



# Long Short Term Memory (LSTM)




The previous hidden state  $h_{t-1}$  and the current input  $x_t$  are used to compute a candidate  $\tilde{c}_t$ :

$$\tilde{c}_t = \tanh(W_c[h_{t-1}, x_t] + b_c)$$

The **cell memory**  $c_t$ , is updated as:

$$c_t = \Gamma_i * \tilde{c}_t + \Gamma_f * c_{t-1}$$

# Long Short Term Memory (LSTM)



The hidden state  $h_t$ , is computed as

$$h_t = \Gamma_o * \tanh(c_t)$$

The **predicted output** is:


$$\hat{y}_t = \sigma_y(W_y h_t + b_y).$$

# Introduction for NLP

Transformer were introduced for NLP as an alternative to sequential models. They rely on **attention mechanisms** to address the lack of long term memory of LSTM.


- Sequences are **computed in parallel**.
- Complexity is in  $\mathcal{O}(T^2 \cdot n)$ .
- Memory is replaced by **attention**.

# Computing in parallel



$$F_{\theta}(x) = P(y_1 \cdots y_T | x_1 \cdots x_T)$$

# Encoder - Decoder structure



The encoder computes a latent vector from the input data, which is fed to the decoder in order to predict the outputs. These sub-networks are trained jointly and are supposed to encourage the Transformer to learn a **meaningful representation** of the data.


# Feed Forward Network

The Feed Forward Network is a combination of linear transformations, with a ReLU activation function:

$$\text{FFN}(x) = \max(0, x \cdot W_1 + b_1) \cdot W_2 + b_2$$


It act as the main **non linearity** of the Transformer.

## Self attention mechanism - intuition



At each time step, what are the part of the entire sequence should we focus on to get the best interpretation ?

# Self attention mechanism - illustration





# Self attention mechanism

$$Q = W_q x \quad (T, q) \quad \text{Queries}$$

$$K = W_k x \quad (T, k) \quad \text{Keys}$$

$$V = W_v x \quad (T, v) \quad \text{Values}$$

$$\text{Scores} = \frac{QK^T}{\sqrt{q}} \quad (T, T)$$

With the softmax function  $\text{softmax}(z) \mapsto \frac{e^{z_i}}{\sum_j e^{z_j}}$  :

## Attention

$$\text{Attention} = \text{softmax}\left(\frac{Q \cdot K}{\sqrt{d_k}}\right)V$$

## Decoder mechanism - single layer example

Transformers-based approaches propose a **model for the conditional distribution of a data  $Y_t$  given past observations.**

**Data representations:** for all  $1 \leq s \leq t$  and all  $1 \leq h \leq n_h$ , define,

$$q^h(s) = W^{h,q} X_s, \quad \kappa^h(s) = W^{h,\kappa} X_s, \quad v^h(s) = W^{h,v} X_s,$$

where  $W^{h,q}$ ,  $W^{h,\kappa}$  and  $W^{h,v}$  are unknown matrices.

**Scores:**

$$\begin{aligned} \text{score}^h(t) &= (q^h(t))^T K^h, \\ \pi^h(t) &= \text{softmax}(\text{score}^h(t) / \sqrt{r}), \end{aligned}$$

where the columns of  $K^h$  are the  $\kappa^h(t-s)$ ,  $1 \leq s \leq t$ .

## Decoder mechanism - single layer example

**Self attention:** for all  $1 \leq s \leq t$ , as,

$$z^h(t) = \sum_{s=1}^t \pi_s^h(t) v^h(t-s),$$

The output is then computed with a Layer normalization step:

$$\begin{aligned} u(t) &= \mathcal{N}(z(t) + X_t), \\ \hat{Y}_t &= \mathcal{N}(\text{FFN}(u(t)) + u(t)), \end{aligned}$$

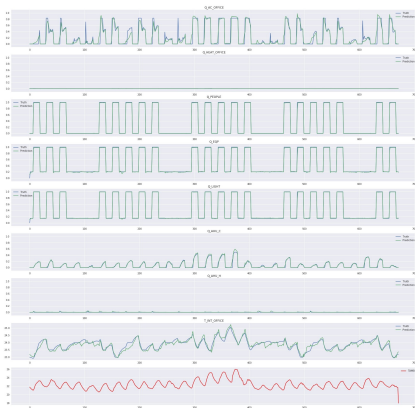
where  $\mathcal{N}$  is a normalizing function.

# Implementation

Our transformer model is implemented using **PyTorch**, the code is available at [github.com/maxjcohen/transformer](https://github.com/maxjcohen/transformer).


Because we do not yet have access to powerful **computing units**, some parameters have been lowered (latent dimension, query/key/value dimension).

## Results: visual interpretation



- Visually, the predictions match the dataset;
- Interior temperature seems much harder to predict;
- Lack of local attention is visible (during the night for example).

## Results: attention maps



We can identify day/night cycles,  
as well as week/weekend.