

MSc Big Data for Business - *MAP 534*

Introduction to machine learning

Supervised classification (II)

Logistic regression & feed forward neural networks

Introduction

Logistic regression

Inference

Feed Forward Neural Networks

Model

Implementation

Setting

- Historical data about **individuals** $i = 1, \dots, n$.
- **Features** vector $X_i \in \mathbb{R}^d$ for each individual i .
- For each i , the individual **belongs to a group** ($Y_i = 0$) or not ($Y_i = 1$).
- $Y_i \in \{0, 1\}$ is the **label** of i .

Objective

- Given a new X (with no corresponding label), **predict a label in $\{0, 1\}$** .
- Use data $\mathcal{D}_n = \{(X_1, Y_1), \dots, (X_n, Y_n)\}$ **to construct a classifier**.

The best solution f^* (which is independent of \mathcal{D}_n) is

$$f^* = \arg \min_{f \in \mathcal{F}} R(f) = \arg \min_{f \in \mathcal{F}} \mathbb{E} [\mathbb{1}_{Y \neq f(X)}] = \arg \min_{f \in \mathcal{F}} \mathbb{P}(Y \neq f(X)).$$

Bayes Predictor (explicit solution)

→ **Binary classification** with 0 – 1 loss:

$$f^*(\mathbf{X}) = \begin{cases} +1 & \text{if } \mathbb{P}\{Y = 1|\mathbf{X}\} \geq \mathbb{P}\{Y = 0|\mathbf{X}\} \\ & \Leftrightarrow \mathbb{P}\{Y = 1|\mathbf{X}\} \geq 1/2, \\ 0 & \text{otherwise.} \end{cases}$$

The explicit solution requires to **know the conditional law of Y given X** ...

How to estimate the conditional law of Y ?

Fully parametric modeling.

Estimate the law of (\mathbf{X}, Y) and use the **Bayes formula** to deduce an estimate of the conditional law of Y : *LDA/QDA, Naive Bayes...*

Parametric conditional modeling.

Estimate the conditional law of Y by a **parametric** law: *linear regression, logistic regression, Feed Forward Neural Networks...*

Nonparametric conditional modeling.

Estimate the conditional law of Y by a **non parametric** estimate: *kernel methods, nearest neighbors...*

The **conditional densities are modeled as multivariate normal**. For all class $k \in \{0, 1\}$, conditionnally on $\{Y = k\}$,

$$X \sim \mathcal{N}(\mu_k, \Sigma_k).$$

Discriminant functions:

$$\psi_k : x \mapsto \ln(g_k(x)) + \ln(\mathbb{P}\{Y = k\}).$$

In a two-classes problem, the optimal classifier is:

$$f^* : x \mapsto \mathbb{1}\{\psi_1(x) > \psi_0(x)\}.$$

QDA (different Σ_k in each class) and LDA ($\Sigma_k = \Sigma$ for all k)

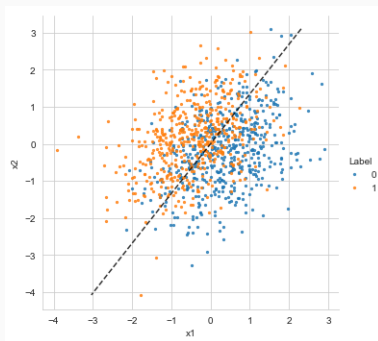
→ May lead to poor results is the **model does not describe the data correctly**.

Fully parametric modeling - Discriminant Analysis

In the LDA case, the classification rule is of the form:

$$f^*(x) = 1 \Leftrightarrow \langle w, x \rangle + b \geq 0,$$

where w and b depends on the model parameters.



→ How to relax the Gaussian assumption ? (**logistic model**).

→ How to design nonlinear classification rules ? (**neural networks**).

Introduction

Logistic regression

Inference

Feed Forward Neural Networks

Model

Implementation

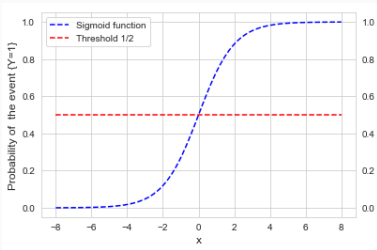
Semi-parametric modelling - logistic regression

- The objective is to **predict the label** $Y \in \{0, 1\}$ based on $X \in \mathbb{R}^d$.
- Logistic regression **models the distribution of Y given X** .

$$\mathbb{P}(Y = 1|X) = \sigma(\langle w, X \rangle + b),$$

where $w \in \mathbb{R}^d$ is a vector of model **weights** and $b \in \mathbb{R}$ is the **intercept**, and where σ is the **sigmoid** function.

$$\sigma : z \mapsto \frac{1}{1 + e^{-z}}.$$



- The sigmoid function is a **model choice to map \mathbb{R} into $(0, 1)$** .
- Another widespread solution for σ is $\sigma : z \mapsto \mathbb{P}(Z \leq z)$ where $Z \sim \mathcal{N}(0, 1)$, which leads to a **probit** regression model.

Log-odd ratio

$$\log \left(\mathbb{P}(Y = 1|X) \right) - \log \left(\mathbb{P}(Y = 0|X) \right) = \langle w, X \rangle + b.$$

Classification rule

Note that

$$\mathbb{P}(Y = 1|X) \geq \mathbb{P}(Y = 0|X)$$

if and only if

$$\langle w, x \rangle + b \geq 0.$$

→ This is a **linear classification** rule.

→ This classifier requires to **estimate w and b** .

Introduction

Logistic regression

Inference

Feed Forward Neural Networks

Model

Implementation

→ $\{(X_i, Y_i)\}_{1 \leq i \leq n}$ are **i.i.d. with the same distribution as (X, Y)** .

Likelihood

$$\begin{aligned}\prod_{i=1}^n \mathbb{P}(Y_i | X_i) &= \prod_{i=1}^n \sigma(\langle w, X_i \rangle + b)^{Y_i} (1 - \sigma(\langle w, X_i \rangle + b))^{1-Y_i}, \\ &= \prod_{i=1}^n \sigma(\langle w, x_i \rangle + b)^{Y_i} \sigma(-\langle w, X_i \rangle - b)^{1-Y_i}\end{aligned}$$

and the **normalized negative loglikelihood** is

$$f(w, b) = \frac{1}{n} \sum_{i=1}^n \log(1 + e^{-Y_i(\langle w, X_i \rangle + b)}) = \frac{1}{n} \sum_{i=1}^n \ell(Y_i, \langle w, X_i \rangle + b).$$

Logistic regression

Compute \hat{w}_n and \hat{b}_n as follows:

$$(\hat{w}_n, \hat{b}_n) \in \operatorname{argmin}_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{n} \sum_{i=1}^n \log(1 + e^{-Y_i(\langle w, X_i \rangle + b)}).$$

→ It is an **average of losses**, one for each sample point.

→ It is a **convex and smooth problem**.

Using the **logistic loss** function

$$\ell : (y, y') \mapsto \log(1 + e^{-yy'})$$

yields

$$(\hat{w}_n, \hat{b}_n) \in \operatorname{argmin}_{w \in \mathbb{R}^d, b \in \mathbb{R}} \frac{1}{n} \sum_{i=1}^n \ell(Y_i, \langle w, X_i \rangle + b).$$

Maximum likelihood estimate

Assume for now that the intercept is 0. Then, the likelihood is,

$$L_n(w) = \prod_{i=1}^n \left(\frac{e^{X_i^T w}}{1 + e^{X_i^T w}} \right)^{Y_i} \left(\frac{1}{1 + e^{X_i^T w}} \right)^{1-Y_i} = \prod_{i=1}^n \left(\frac{e^{X_i^T w Y_i}}{1 + e^{X_i^T w}} \right).$$

And the **negative log-likelihood** is

$$\ell_n(w) = -\log(L_n(w)) = \sum_{i=1}^n \left(-Y_i X_i^T w + \log(1 + e^{X_i^T w}) \right).$$

Derivatives

$$\begin{aligned} \frac{\partial (\log(L_n(w)))}{\partial w_j} &= \sum_{i=1}^n \left(Y_i X_{ij} - \frac{x_{ij} e^{X_i^T w}}{(1 + e^{X_i^T w})} \right) \\ &= \sum_{i=1}^n X_{ij} (Y_i - \sigma(\langle w, X_i \rangle)). \end{aligned}$$

→ **No explicit solution** for the maximizer of the loglikelihood... Parameter estimate obtained using **gradient based optimization** (see next lesson).

Maximum likelihood estimate

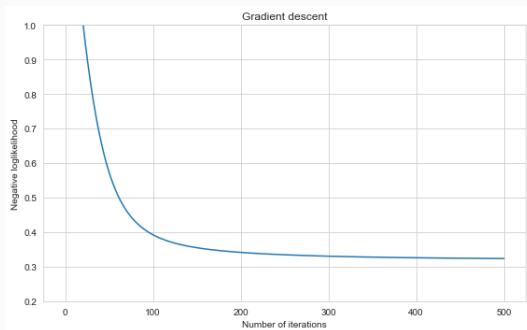
The negative loglikelihood

$$\ell_n(w) = -\log(L_n(w)) = \sum_{i=1}^n \left(-Y_i X_i^T w + \log(1 + e^{X_i^T w}) \right).$$

is minimized using a gradient descent algorithm.

Starting with an **initial estimate** $w^{(0)}$, for all $k \geq 1$, set

$$w^{(k)} = w^{(k-1)} - \eta_k \nabla \ell_n(w^{(k-1)}).$$

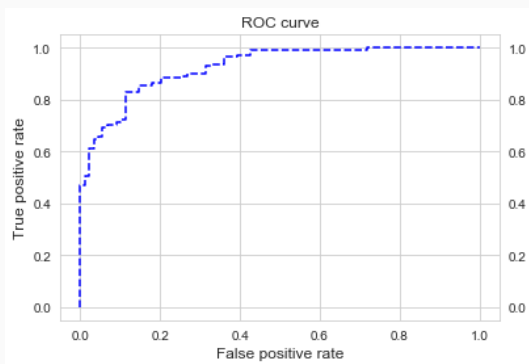


Maximum likelihood estimate

Let (w^*, b^*) be the parameter estimates **after the gradient descent algorithm**.

The usual logistic regression classifier is $f^*(X) = 1 \Leftrightarrow \mathbb{P}(Y = 1|X) > 1/2$.

Sensitivity of the classifier to this threshold: for each value $p^* \in (0, 1)$ the ROC curve classifies individuals using $f^*(X) = 1 \Leftrightarrow \mathbb{P}(Y = 1|X) > p^*$ and plots the True positive rate as a function of the False positive rate.



The **gradient of the negative loglikelihood** is,

$$\nabla \ell_n(w) = - \sum_{i=1}^n Y_i X_i + \sum_{i=1}^n \frac{\exp(\langle X_i, w \rangle)}{1 + \exp(\langle X_i, w \rangle)} X_i .$$

On the other hand, for all $1 \leq i \leq n$ and all $1 \leq j \leq d$,

$$\partial_j \left(\frac{\exp(\langle X_i, w \rangle)}{1 + \exp(\langle X_i, w \rangle)} X_i \right) = \frac{\exp(\langle X_i, w \rangle)}{(1 + \exp(\langle X_i, w \rangle))^2} X_{ij} X_i ,$$

where X_{ij} is the j th component of X_i .

Then, the **Hessian matrix** is

$$(H_n(w))_{\ell j} = \sum_{i=1}^n \frac{\exp(\langle X_i, w \rangle)}{(1 + \exp(\langle X_i, w \rangle))^2} X_{ij} X_{i\ell} ,$$

that is,

$$H_n(w) = \sum_{i=1}^n \frac{\exp(\langle X_i, w \rangle)}{(1 + \exp(\langle X_i, w \rangle))^2} X_i X_i^T .$$

$H_n(\beta)$ is a semi positive definite matrix, which implies that $\ell_n(\beta)$ is convex.

Asymptotic properties

Assumptions

→ $\hat{w}_n \rightarrow w^*$ **almost surely**.

→ There exists a continuous and nonsingular function H such that $n^{-1}H_n(w)$ **converges to $H(w)$** , uniformly in a ball around w^* .

For all $t \in \mathbb{R}^d$, using a Taylor expansion,

$$\mathbb{E} \left[\exp \left(-\frac{1}{\sqrt{n}} \langle t, \nabla \ell_n(w^*) \rangle \right) \right] \rightarrow_{n \rightarrow \infty} \exp \left(\frac{1}{2} t^T H(w^*) t \right).$$

Therefore,

$$-\nabla \ell_n(w^*) / \sqrt{n} \Rightarrow \mathcal{N}(0, H(w^*)).$$

On the other hand, by **Slutsky lemma**,

$$\sqrt{n}(\hat{w}_n - w^*) \Rightarrow \mathcal{N}(0, H(w^*)^{-1}).$$

Confidence interval

$\rightarrow \sqrt{n}(\hat{w}_j - w_j^*)$ converges in distribution to a centered Gaussian random variable with variance $(H(w^*)^{-1})_{jj}$.

Almost surely,

$$\hat{\sigma}_{n,j}^2 = (nH_n(\hat{w}_n)^{-1})_{jj} \rightarrow_{n \rightarrow \infty} (H(w^*)^{-1})_{jj}.$$

Then,

$$\sqrt{\frac{n}{\hat{\sigma}_{n,j}^2}}(\hat{w}_{n,j} - \beta_j^*) \rightarrow_{n \rightarrow \infty} \mathcal{N}(0, 1).$$

An asymptotic confidence interval $\mathcal{I}_{n,\alpha}$ of level $1 - \alpha$ is then

$$\mathcal{I}_{n,\alpha} = \left[\hat{w}_{n,j} - z_{1-\alpha/2} \sqrt{\frac{\hat{\sigma}_{n,j}^2}{n}}, \hat{\beta}_{n,j} + z_{1-\alpha/2} \sqrt{\frac{\hat{\sigma}_{n,j}^2}{n}} \right],$$

where $z_{1-\alpha/2}$ is the quantile of order $1 - \alpha/2$ of $\mathcal{N}(0, 1)$.

Introduction

Logistic regression

Inference

Feed Forward Neural Networks

Model

Implementation

Introduction

Logistic regression

Inference

Feed Forward Neural Networks

Model

Implementation

Softmax function

- The objective is to **predict the label** $Y \in \{1, \dots, M\}$ based on $X \in \mathbb{R}^d$.
- Softmax regression **models the distribution of Y given X** .

The model

For all $1 \leq m \leq M$,

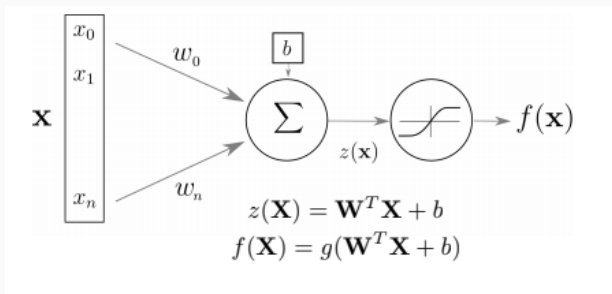
$$z_m = \langle w_m, X \rangle + b_m,$$

$$\mathbb{P}(Y = m|X) = \text{softmax}(z)_m,$$

where $z \in \mathbb{R}^M$, $w_m \in \mathbb{R}^d$ is a vector of model **weights** and $b_m \in \mathbb{R}$ is an **intercept**, and where softmax is the **softmax** function: for all $1 \leq m \leq M$,

$$\text{softmax}(z)_m = \frac{\exp(z_m)}{\sum_{j=1}^M \exp(z_j)}.$$

A neuron



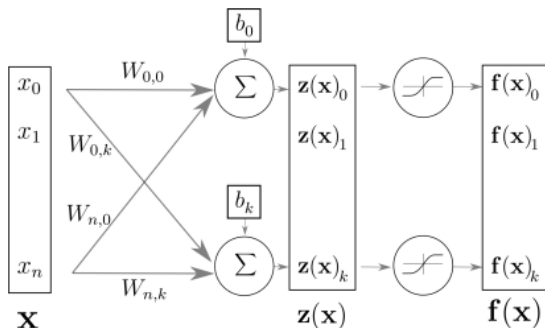
→ \mathbf{X} **input** in \mathbb{R}^d .

→ $z(\mathbf{X})$ **pre-activation** in \mathbb{R}^M , with **weight** $\mathbf{W} \in \mathbb{R}^{d \times M}$ and **bias** $b \in \mathbb{R}^M$.

→ g **softmax function**.

One neuron is a multi-class extension of the logistic regression model.

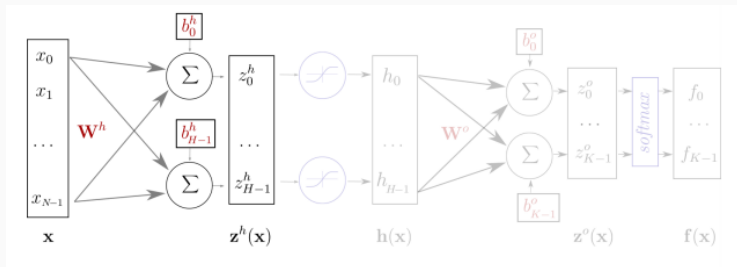
Layer of neurons and hidden states



$$f(\mathbf{X}) = g(\mathbf{z}(\mathbf{X})) = g(\mathbf{W}\mathbf{X} + \mathbf{b})$$

- \mathbf{X} **input** in \mathbb{R}^d .
- $\mathbf{z}(\mathbf{X})$ **pre-activation** in \mathbb{R}^k , with **weight** $\mathbf{W} \in \mathbb{R}^{d \times k}$ and **bias** $\mathbf{b} \in \mathbb{R}^k$.
- g **any activation function** (nonlinear & nondecreasing function).
- $\mathbf{f}(\mathbf{X})$ **hidden state** in \mathbb{R}^k which may be used as input of a new neuron...

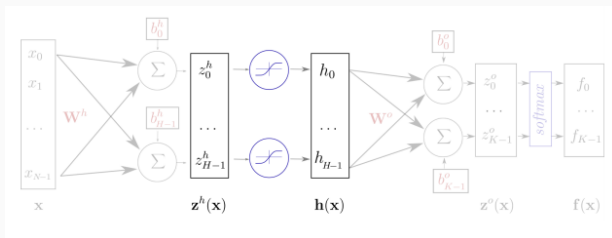
Feed Forward Network



→ \mathbf{X} **input** in \mathbb{R}^d .

→ $\mathbf{z}^h(\mathbf{X})$ **pre-activation** in \mathbb{R}^H , with **weight** $\mathbf{W}^h \in \mathbb{R}^{d \times H}$ and **bias** $\mathbf{b}^h \in \mathbb{R}^H$.

Feed Forward Network

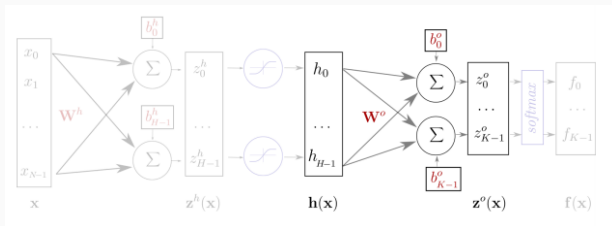


→ X input in \mathbb{R}^d .

→ $z^h(X)$ pre-activation in \mathbb{R}^H , with weight $W^h \in \mathbb{R}^{d \times H}$ and bias $b^h \in \mathbb{R}^H$.

→ g any activation function to produce $h \in \mathbb{R}^H$.

Feed Forward Network



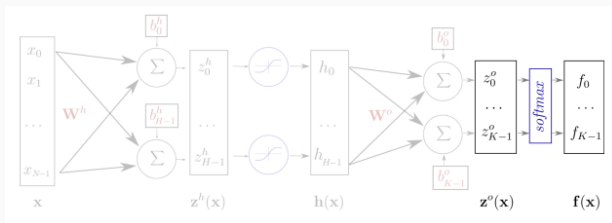
→ \mathbf{X} **input** in \mathbb{R}^d .

→ $\mathbf{z}^h(\mathbf{X})$ **pre-activation** in \mathbb{R}^H , with **weight** $W^h \in \mathbb{R}^{d \times H}$ and **bias** $b^h \in \mathbb{R}^H$.

→ g **any activation function** to produce $\mathbf{h} \in \mathbb{R}^H$.

→ $\mathbf{z}^o(\mathbf{X})$ **pre-activation** in \mathbb{R}^M , with **weight** $W^o \in \mathbb{R}^{H \times M}$ and **bias** $b^o \in \mathbb{R}^M$.

Feed Forward Network



→ X **input** in \mathbb{R}^d .

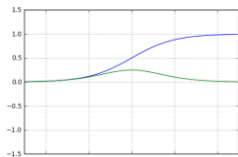
→ $z^h(X)$ **pre-activation** in \mathbb{R}^H , with **weight** $W^h \in \mathbb{R}^{d \times H}$ and **bias** $b^h \in \mathbb{R}^H$.

→ g **any activation function** to produce $h \in \mathbb{R}^H$.

→ $z^o(X)$ **pre-activation** in \mathbb{R}^M , with **weight** $W^o \in \mathbb{R}^{H \times M}$ and **bias** $b^o \in \mathbb{R}^M$.

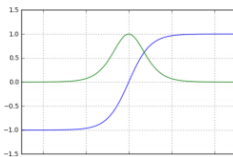
→ Apply the **softmax function to produce the output**, i.e. $\mathbb{P}(Y = m|X)$ for $1 \leq m \leq M$.

Activation functions



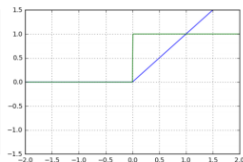
$$\text{sigm}(x) = \frac{1}{1 + e^{-x}}$$

$$\text{sigm}'(x) = \text{sigm}(x)(1 - \text{sigm}(x))$$



$$\tanh(x) = \frac{e^{2x} - 1}{e^{2x} + 1}$$

$$\tanh'(x) = 1 - \tanh(x)^2$$



$$\text{relu}(x) = \max(0, x)$$

$$\text{relu}'(x) = 1_{x>0}$$

→ As there is no modelling assumptions anymore, **virtually any activation function** may be used.

→ The rectified linear unit (ReLU) activation function $\sigma(x) = \max(0, x)$ and its extensions are the default recommendation in modern implementations (Jarrett et al., 2009; Nair and Hinton, 2010; Glorot et al., 2011a), (Maas et al., 2013), (He et al., 2015). One of the major motivations arise from the **gradient based parameter optimization which is numerically more stable with this choice**.

Introduction

Logistic regression

Inference

Feed Forward Neural Networks

Model

Implementation

→ This dataset contains images representing handwritten digits. Each image is made of **28 x 28 pixels**, and each pixel is represented by an integer (gray level). These arrays can be flattened into vectors in \mathbb{R}^{784} .

→ The **labels in $\{0, \dots, 9\}$ are represented using one-hot-encoding** and grayscale of each pixel in $\{0, \dots, 255\}$ are normalized to be in $(0, 1)$.

```
from keras.datasets import mnist
# Number of classes
num_classes = 10
# input image dimensions
img_rows, img_cols = 28, 28

# the data, shuffled and split between train and test sets
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')

print('x_train shape:', x_train.shape)
print('x_test shape:', x_test.shape)
print('y_train shape:', y_train.shape)
print('y_test shape:', y_test.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')
```

```
x_train shape: (60000, 28, 28, 1)
x_test shape: (10000, 28, 28, 1)
y_train shape: (60000,)
y_test shape: (10000,)
60000 train samples
10000 test samples
```

- This dataset contains images representing handwritten digits. Each image is made of **28 x 28 pixels**, and each pixel is represented by an integer (gray level). These arrays can be flattened into vectors in \mathbb{R}^{784} .
- The **labels in $\{0, \dots, 9\}$ are represented using one-hot-encoding** and grayscale of each pixel in $\{0, \dots, 255\}$ are normalized to be in $(0, 1)$.

```
plt.figure(figsize=(8, 2))
for i in range(4):
    plt.subplot(1, 4, i+1)
    plt.imshow(x_train[i].reshape(28, 28),
               interpolation="none", cmap="gray_r")
    plt.title('Label = %d' % y_train[i], fontsize=14)
    plt.axis("off")
plt.tight_layout()
```

Label = 5

5

Label = 0

0

Label = 4

4

Label = 1

1

The model with Keras

```
model_ffnn = Sequential()

model_ffnn.add(Flatten(input_shape=input_shape))

model_ffnn.add(Dense(128, activation='relu'))

model_ffnn.add(Dense(num_classes, activation='softmax'))
```

```
model_ffnn.compile(
    loss=keras.losses.categorical_crossentropy,
    optimizer=keras.optimizers.Adagrad(),
    metrics=['accuracy']
)

model_ffnn.summary()
```

| Layer (type) | Output Shape | Param # |
|---------------------------|--------------|---------|
| flatten_1 (Flatten) | (None, 784) | 0 |
| dense_1 (Dense) | (None, 128) | 100480 |
| dense_2 (Dense) | (None, 10) | 1290 |
| Total params: 101,770 | | |
| Trainable params: 101,770 | | |
| Non-trainable params: 0 | | |

Figure 1: Feed Forward Neural network. h_1 is obtained with the **RELU activation function and is in \mathbb{R}^{128}** . The last layer is $h_2 \in \mathbb{R}^{10}$ and is **obtained with the softmax activation function** so that each component m models $\mathbb{P}(Y = m|X)$. This neural network with one hidden layer relies on 101.770 parameters.

→ This model relies on more than **100.000 unknown parameters** which should be estimated.

→ As for the logistic regression and the discriminant analysis, a common choice is to **minimize the negative loglikelihood of the data**:

$$\theta \mapsto -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^{10} \mathbb{1}_{Y_i=k} \log \mathbb{P}_{\theta}(Y_i = k | X_i).$$

→ The negative loglikelihood is computed using $n = 60.000$ training samples and **minimized using gradient descent algorithms** - see next lesson.

→ Then, the performance of the model is assessed using 10.000 new (test) samples: the **accuracy is the frequency of labels which are well predicted by the model with the estimated parameters**.

```

batch_size = 32
epochs = 8

# Run the train
history = model_ffnn.fit(x_train, y_train,
                        batch_size=batch_size,
                        epochs=epochs,
                        verbose=1,
                        validation_data=(x_test, y_test))
score = model_ffnn.evaluate(x_test, y_test, verbose=0)
print('Test loss:', score[0])
print('Test accuracy:', score[1])

```

```

plt.figure(figsize=(5, 4))
plt.plot(history.epoch, history.history['acc'], lw = 1, label='Training')
plt.plot(history.epoch, history.history['val_acc'], lw = 1, label='Testing')
plt.legend()
plt.title('Accuracy of softmax regression', fontsize=16)
plt.xlabel('Epoch', fontsize=14)
plt.ylabel('Accuracy', fontsize=14)
plt.tick_params(labelright=True)
plt.grid('True')
plt.tight_layout()

```

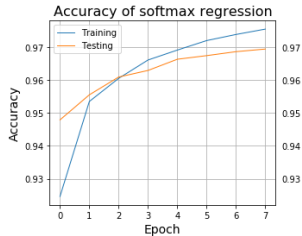


Figure 2: Minimization of the negative loglikelihood using a gradient descent algorithm (here AdaGrad). The gradient is computed using batches of 32 observations and the whole data set is used 8 times.